

# **I/O CONTROLLER INTERFACE AND I/O DEVICES ON THE LC4**

# I/O Devices and Controllers

Most I/O devices are not purely digital themselves ...

- Electro-mechanical: e.g. keyboard, mouse, disk, motor
- Analog/digital: e.g. touchscreen, network interface, monitor, speaker, mic

... all have digital interfaces presented by **I/O Controller**

- *I/O Device (analog/digital mix) talks to controller*
- CPU (digital) talks to controller
- *Controller acts as a translator: digital (CPU) <-> analog (device)*



# I/O Controller to CPU Interface

I/O controller interface abstracts I/O device as “device registers”

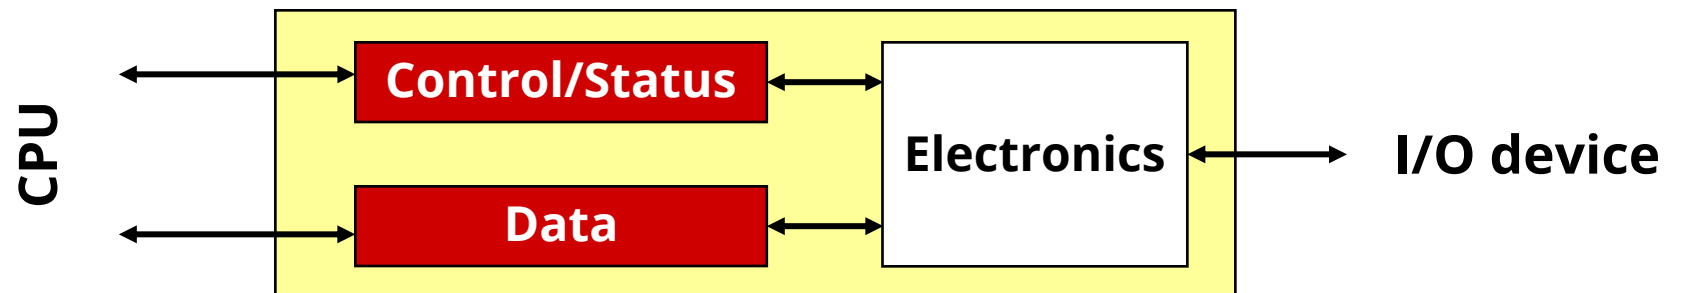
- **Control/Status**: may be one register or two
- **Data**: may be more than one of these

For input devices

- CPU checks **status register** if input is available
- Reads input from **data register** (or waits if no input available)

For output devices

- CPU checks **status register** to see if it can write
- Writes output to **data register**



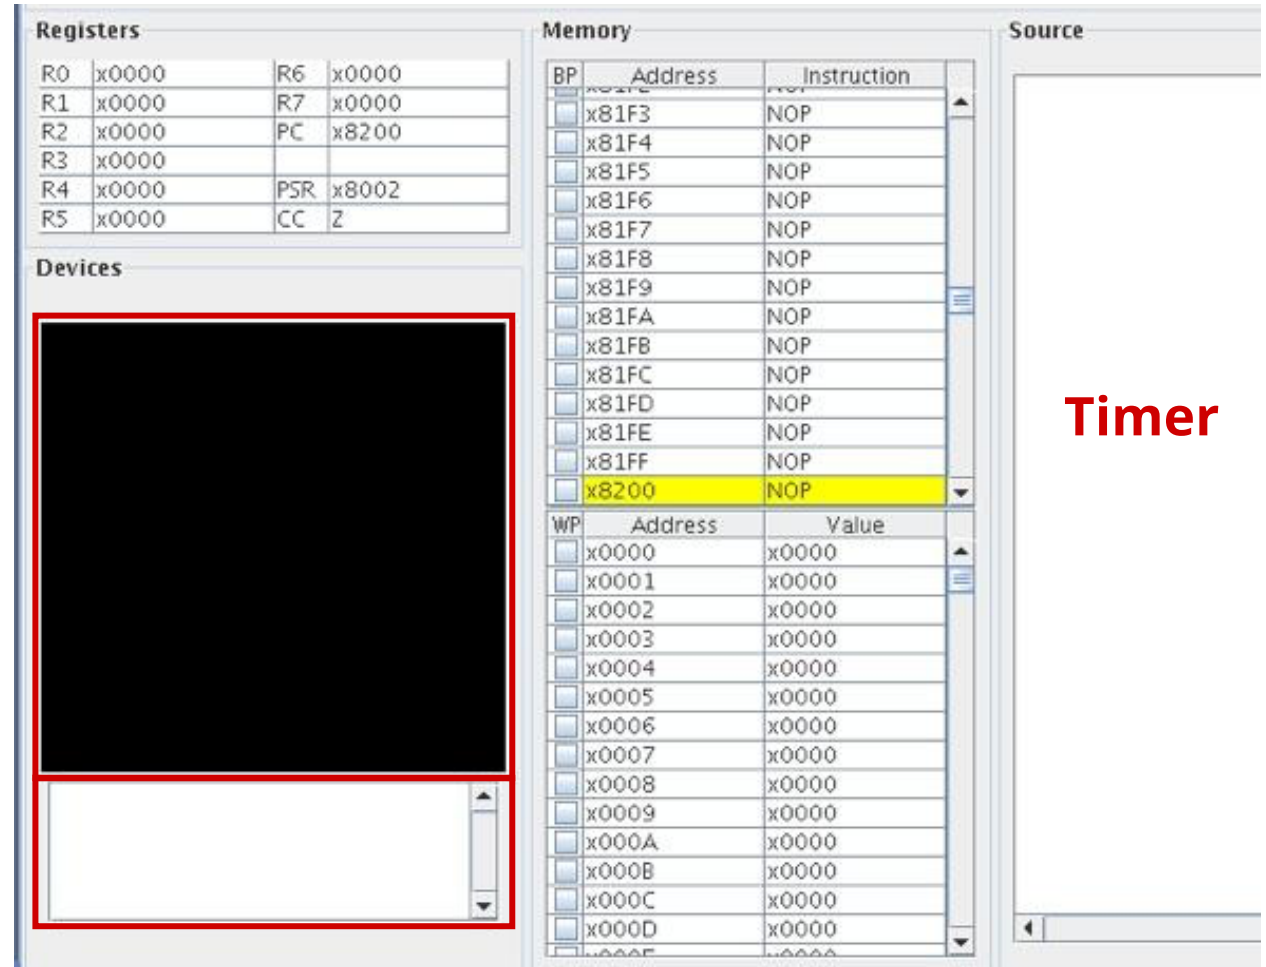
# LC4 I/O Devices

## LC4 has 4 I/O devices

- Keyboard (input)
- ASCII console (output)
- 128x124 16-bit RGB pixel display (output)
- Timer (not really an I/O device but looks like one to software)

**Video display**

**Keyboard/console**



**Timer**

# HOW ARE DEVICE REGISTER READS/ WRITES PERFORMED?

# How are Device Register Reads/Writes Performed?

## Two options (aren't there always?)

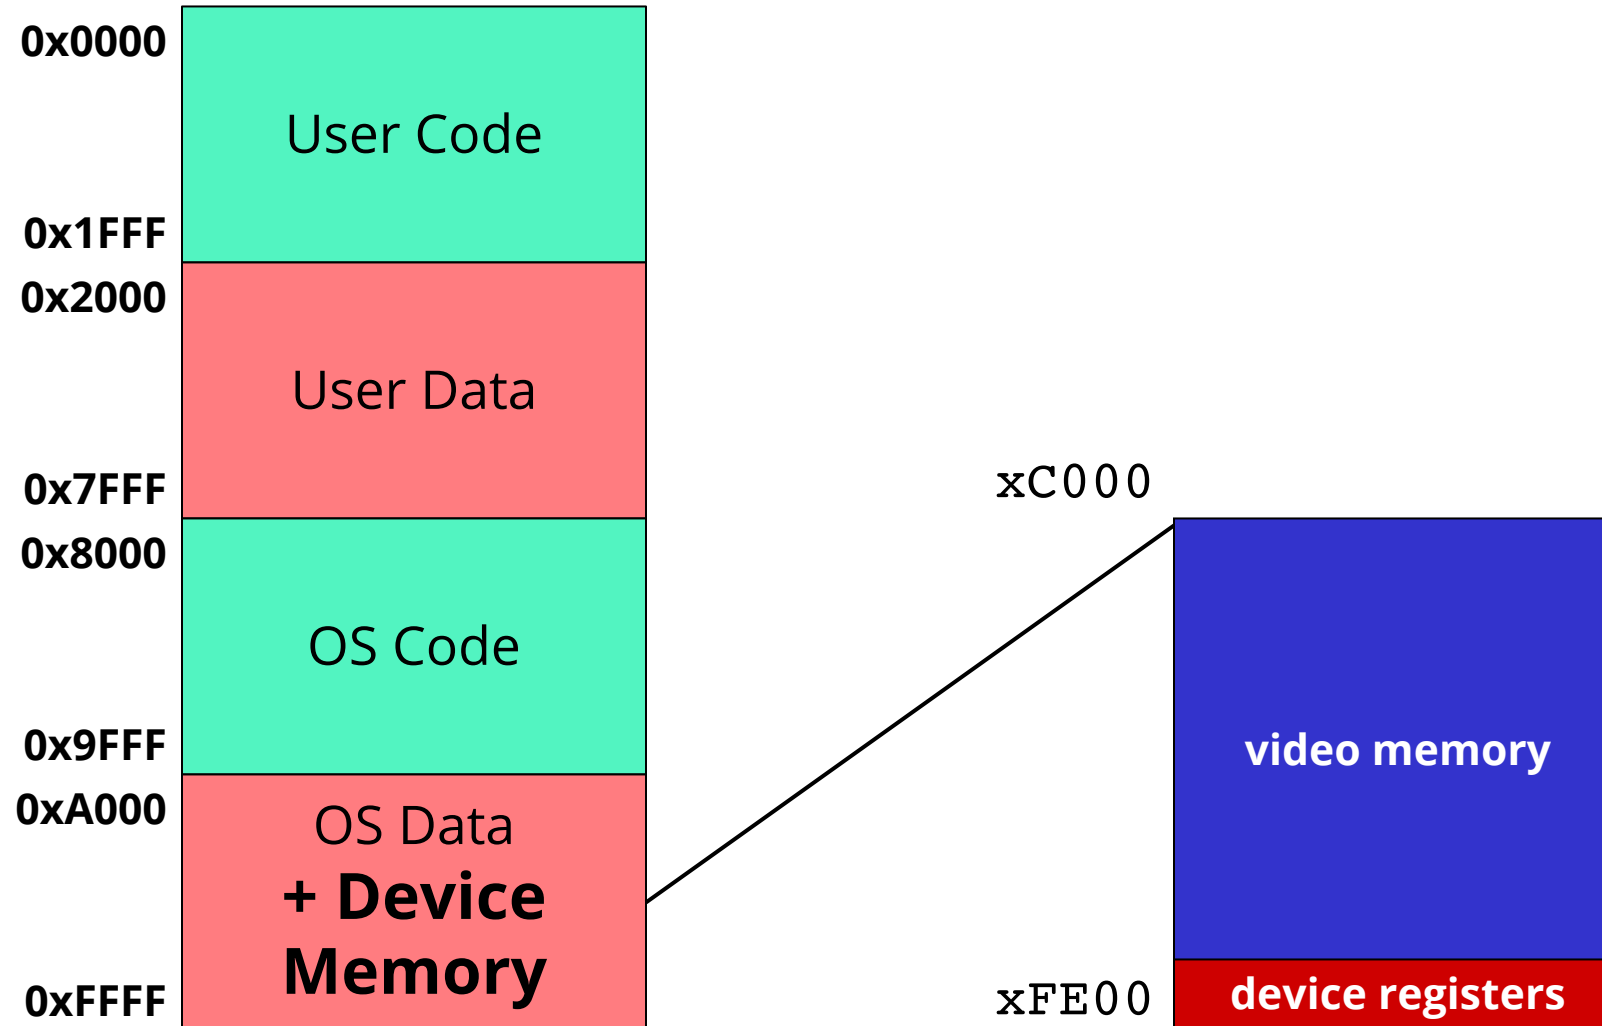
We could create new "I/O instructions" for the ISA

- Designate opcode(s) for I/O
- Register and operation encoded in instruction

Memory-mapped I/O (*we already have these in ISA: LDR/STR*)

- Assign a memory address to each device register
- Use conventional loads and stores
- Hardware intercepts loads/stores to these address
- No actual memory access performed
- LC4 (and most other platforms) do this

# LC4 Data Memory (The Other Half)



# LC4 ASCII I/O Device Registers

Keyboard status register (**KBSR**): xFE00

- KBSR[15] is 1 if keyboard has new character

Keyboard data register (**KBDR**): xFE02

- KBDR[7:0] is last character input on keyboard

ASCII display status register (**ADSR**): xFE04

- ADSR[15] is 1 if console ready to display next character

ASCII display data register (**ADDR**): xFE06

- ADDR[7:0] is written to console

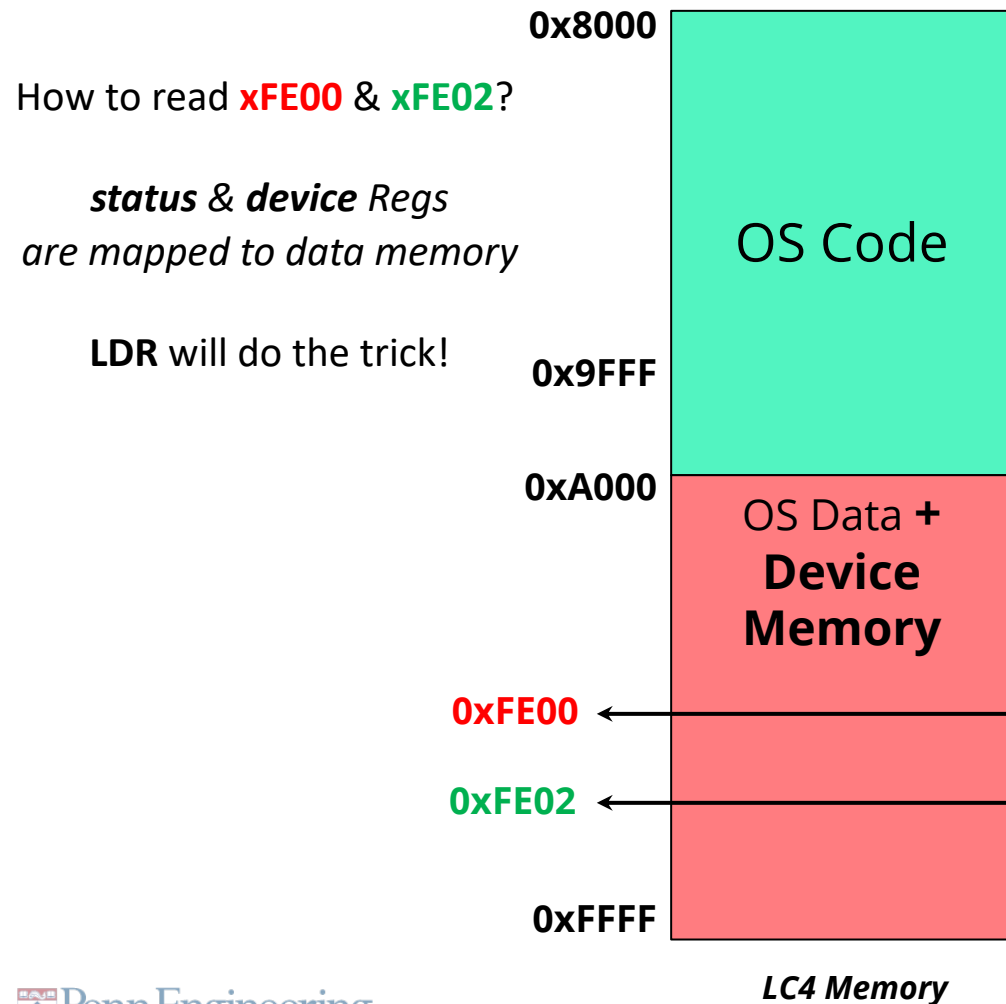
xFE00

device registers

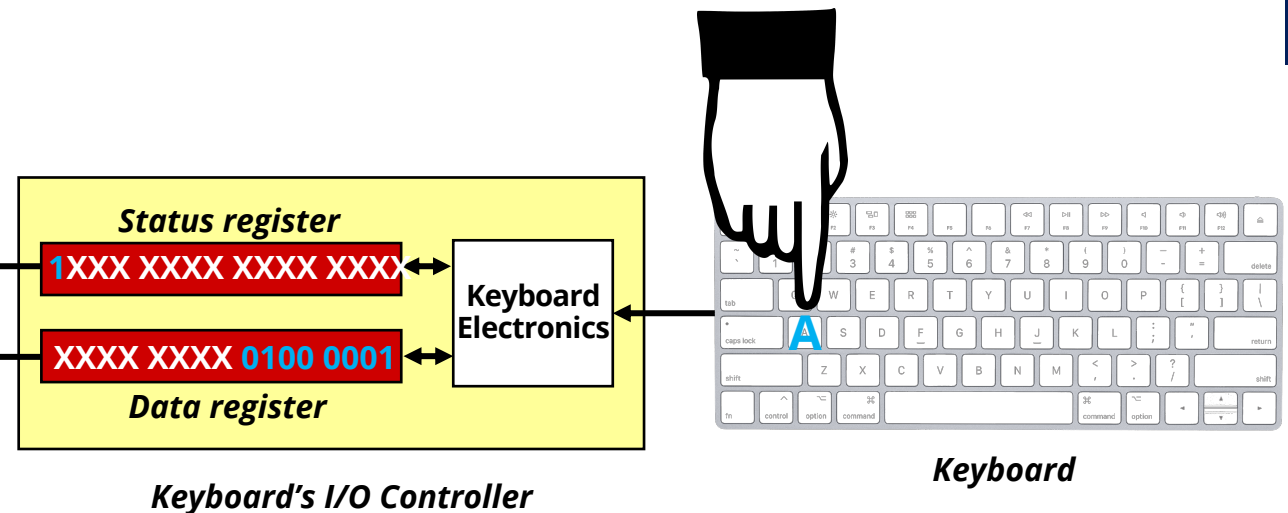


# MEMORY “MAPPED” I/O – HOW DOES IT WORK?

# Memory “Mapped” I/O – How Does It Work:



1. User presses 'A' key on keyboard
2. Electronics box converts 'A' to ASCII: **0100 0001**
  - *places this data into **data register***
3. Electronics box sets MSB of **status register** to '1'
4. Programmer checks **status register**
  - *by reading address **0xFE00***
5. If MSB = 1, programmer reads **data register**
  - *by reading address **0xFE02***
6. Electronics box reset **status register** after read



# DEVICE 1: KEYBOARD – READING INPUT

# Example 1: Reading from Keyboard

*; code will read 1 character from the keyboard, store it in R0*

```
OS_KBSR_ADDR .UCONST xFE00 ; 'alias' for keyboard status reg
OS_KBDR_ADDR .UCONST xFE02 ; 'alias' for keyboard data reg
```

.CODE

```
GETC ; a LABEL for now (perhaps subroutine someday)
    LC R0, OS_KBSR_ADDR ; R0 = address of keyboard status reg
    LDR R0, R0, #0 ; R0 = value of keyboard status reg
    BRzp GETC ; if R0[15]=1, data is waiting!
                ; else, loop and check again...
```

*;; reaching here, means data is waiting in keyboard data reg*

```
LC R0, OS_KBDR_ADDR ; R0 = address of keyboard data reg
LDR R0, R0, #0 ; R0 = value of keyboard data reg
```

***When complete, R0 contains ASCII  
character from keyboard***

# Assembler Constants Using Directives

```
OS_KBSR_ADDR .UCONST xFE00 ; 'alias' for keyboard status reg
```

## .UCONST

- Recall, this is an assembly “directive”
- Mnemonic: `.UCONST UIMM16`
- Function: associate UIMM16 with preceding label
- Defines an *unsigned* 16-bit constant (unlike `.CONST`)
- Why not just use `.FILL`?
  - `.FILL` directives show up in data memory
  - `.UCONST` directives don't

# Assembler Constants Using Directives

```
OS_KBSR_ADDR .UCONST xFE00 ; 'alias' for keyboard status reg  
LC R0, OS_KBSR_ADDR ; R0 = address of keyb status reg
```

LC (Load Constant)

- Assembler pseudo-instruction similar to LEA
- Expands into CONST, HICONST pair
- Loads value *at* label rather than address *of* label
  - ***LEA reads address of the label***

# DEVICE 2: ASCII DISPLAY – WRITING OUTPUT

## Example 2: Read from Keyboard, Print to Screen

; reads 1 character from the keyboard, prints it to ASCII display

```
OS_KBSR_ADDR .UCONST xFE00
OS_KBDR_ADDR .UCONST xFE02
OS_ADSR_ADDR .UCONST xFE04
OS_ADDR_ADDR .UCONST xFE06
```

}

Aliases for keyboard status & data regs

}

Aliases for **ASCII** display status & data regs

.CODE

GETC

```
LC R0, OS_KBSR_ADDR ;; loop while KBSR[15]==0
```

```
LDR R0, R0, #0
```

```
BRzp GETC
```

```
LC R0, OS_KBDR_ADDR
```

```
LDR R0, R0, #0 ;; read data from keyboard
```

PUTC

```
LC R1, OS_ADSR_ADDR ;; loop while ADSR[15]==0
```

```
LDR R1, R1, #0
```

```
BRzp PUTC
```

```
LC R1, OS_ADDR_ADDR
```

```
STR R0, R1, #0 ;; write R0 to ASCII display
```

}

Get character  
from keyboard

}

Print character  
to ASCII display




# DEVICE 3 & 4: TIMER & VIDEO

# LC4 I/O Device Registers: Timer / Video

## TIMER:

- Timer interval register (**TIR**): xFE0A
  - Set desired time in TIR (in msec)
- Timer status register (**TSR**): xFE08
  - TSR[15] is 1 if timer has “gone off”, sets itself to 0 after read
- Works like an egg timer, set desired time in TIR,
  - Then poll/check TSR to see if time has expired

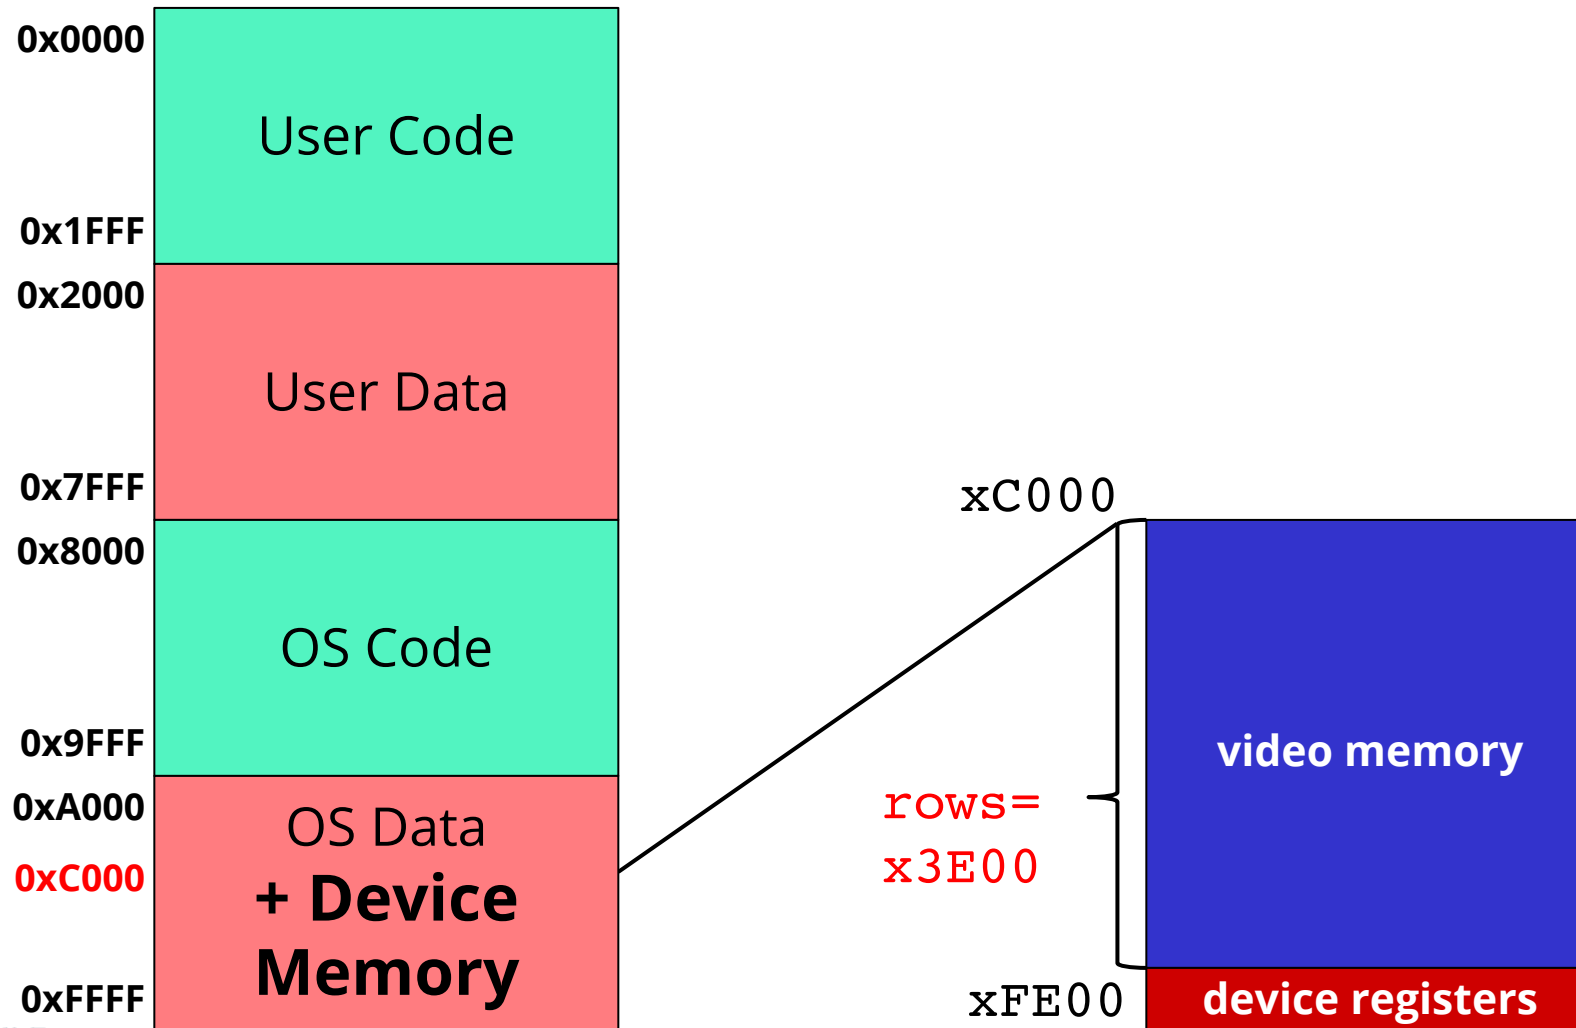


## VIDEO:

- Video display **control** register (**VDCR**): xFE0C
- Video display's many **data** registers: xC000–xFDFF
  - *We'll talk about how to work with video next!*

# DEVICE 4: VIDEO – HOW TO CALCULATE PIXELS

# Where Does Video Memory Live?




# LC4 Pixel-Based Video Display

LC4 has a 128x124 16b RGB (32K color) pixel display

- 128 columns (0-127) and 124 rows (0-123)

Entire display is memory-mapped

- One memory location per pixel
- Memory region `xC000-xFFFF`
  - *`xC000-xC07F` is first row, `xC080-xC0FF` is second row, etc.*
- Write to memory location to set pixel color
- Your job: compute location of pixel
- Then STR color to that address




# Addressing a Pixel

What's the difficult thing here?

- Computing memory address that corresponds to `vmem[y][x]`

```
.ADDR xC000
OS_VIDEO_MEM      .BLKW x3E00    ; why 3E00?
OS_VIDEO_NUM_COLS .UCONST #128
OS_VIDEO_NUM_ROWS .UCONST #124
```

- Logically 2D, but 1D in memory
- Row-major** order (`vmem[y][x]`)
  - `vmem[y][x]` – pixel on row `y`, col `x`
- Pixel at `vmem[2][5]` stored at
  - `xC000 + (2 * 128) + 5`
- In general `vmem[y][x]` stored at
  - `xC000 + (y * 128) + x`
- Note indexing from upper left corner of the display



# USER MEMORY VS. OS MEMORY

# LC-4 Memory Map

Recall the full LC-4 Memory Map

- We have 2 Program Memories
- And 2 Data Memories


## User Region

- Programs run by users (e.g.: factorial program)
- Processes run in user mode have **PSR[15]=0**
- **NOT allowed** to access OS locations in memory

## Operating System Region

- Programs run by OS (e.g.: I/O device programs)
- Processes run in OS mode have **PSR[15]=1**
- **Allowed** to access OS & User locations in memory

## Processor Status Register





# Installing GETC in OS Program Memory

*; subroutine to read 1 character from the keyboard, return it in R0*

```
OS_KBSR_ADDR .UCONST xFE00
OS_KBDR_ADDR .UCONST xFE02
```

```
.CODE
.ADDR x0000
GETC
    LC R0, OS_KBSR_ADDR
    LDR R0, R0, #0
    BRzp GETC

    LC R0, OS_KBDR_ADDR
    LDR R0, R0, #0
```

- There is a slight problem with this code
  - Since it will be loaded into program memory: **x0000**
    - *the LDR statements will fail!*
  - Programs running in USER program memory:
    - *have PSR[15]=0*
    - *they cannot access OS data memory*

# Installing GETC in OS Program Memory

; subroutine to read 1 character from the keyboard, return it in R0

```
OS_KBSR_ADDR .UCONST xFE00
OS_KBDR_ADDR .UCONST xFE02
```

```
.OS
.CODE
.ADDR x8000
SUB_GETC
    LC R0, OS_KBSR_ADDR
    LDR R0, R0, #0
    BRzp GETC

    LC R0, OS_KBDR_ADDR
    LDR R0, R0, #0
    RET
```

- These 3 red bolded directives:
  - **.OS .CODE .ADDR x8000**
  - instruct the assembler, to tell the loader, to load this program into OS program memory
- When the LC4 executes this code, PSR[15]=1
  - *since the PSR[15]=1,*
  - *this program will be allowed to LDR from OS data memory*
- We have one small problem...
  - *what if we turn this into a subroutine*
  - *how can we call this subroutine from user space?*

# Calling GETC from User Program Memory

When a program is running in User Program Memory,  $\text{PSR}[15] = 0$

- We know we can't LDR/STR to device memory
- This is why we installed our GETC subroutine in OS program memory, so that  $\text{PSR}[15]=1$
- But what if we are running a program in USER program memory ( $\text{PSR}[15]=0$ )...
  - *and we want to JSR to a subroutine in OS program memory (to perform some I/O)?*
  - JSR won't allow it! Recall JSR semantics:  $\text{PC} = (\text{PC} \ \& \ 0x8000) \mid (\text{IMM11} \ll 4)$
  - JMP won't allow it either...
    - Why?? The LC4 will kill any program that attempts to jump into OS memory, since its  $\text{PSR}[15]=0$
- What can we do?
  - There 1 instruction in the ISA that can change  $\text{PSR}[15]$  from 0 to 1...do you recall it?
    - **The TRAP instruction!**

# WHAT IS A TRAP?

# TRAP Instruction vs. JSR Instruction

Mnemonic	Semantics	Encoding
TRAP UIMM8	R7 = PC+1, PC = (x8000   UIMM8), <b>PSR[15] = 1</b>	<b>1111</b> ----UUUUUUUU
JSR IMM11	R7 = PC+1, PC = (PC&x8000)   (IMM11<<4)	<b>01101</b> IIIIIIIIIIII

The TRAP instruction is very similar to a JSR:

- It saves PC+1 into R7
- It updates the PC to an offset you specify
- **BUT, it also elevates the **privilege level** of the CPU from 0 to 1**

The purpose of the TRAP instruction:

- Allow a program running in USER Program Memory,
  - to call a subroutine installed in OS Program Memory
- *Fancy name for subroutines in OS? You guessed it: TRAPS*

# RTI Instruction vs. RET Pseudo-Instruction

Mnemonic	Semantics	Encoding
RTI	PC = R7, PSR[15] = 0	1000-----
RET	JMP R7, which simply sets: PC = R7	11000--111-----

The RTI instruction is very similar to a RET:

- It restores the PC back to the value saved in R7 (just like RET)
- BUT, it also lowers the **privilege level** of the CPU from 1 to 0

The purpose of the RET instruction:

- Allow a subroutine running in the OS program memory
  - to return back to a caller in the USER program memory

# Installing GETC in OS Program Memory

```
; User Program Memory
```

```
.CODE
```

```
.ADDR x0000
```

```
; doing some fun stuff, like computing factorials!
```

```
; now, let's get a character from the keyboard!
```

```
;
```

```
TRAP x00 ; saves R7=PC+1, sets PC = x8000 | x00, and PSR[15]=1
```

```
; upon return, do something with R0
```

```
; OS Program Memory
```

```
.OS
```

```
.CODE
```

```
.ADDR x8000
```

```
SUB_GETC
```

```
    LC R0, OS_KBR_ADDR
```

```
    LDR R0, R0, #0
```

```
    BRzp GETC
```

```
    LC R0, OS_KBR_ADDR
```

```
    LDR R0, R0, #0
```

```
RTI
```

```
; gets character from keyboard, saves it in R0
```

```
; sets PC = R7 and restores PSR[15]=0
```

# User Mode vs OS Mode – An example

## User Mode



**Mild mannered reporter**

**Decent typist**

***PSR[15]=0...No powers***

## SuperUser Mode



**All powerful!**

***PSR[15]=1***



# THE TRAP VECTOR TABLE

# The Limits of the TRAP Instruction

Mnemonic	Semantics	Encoding
TRAP UIMM8	$R7 = PC + 1,$ $PC = (x8000 \mid UIMM8),$ $PSR[15] = 1$	1111----UUUUUUUU

The TRAP instruction is limited:

- It does not allow the user to jump just anywhere in the OS
- In only let's the user jump into the first 256 rows of the OS
- **Why? So that the user doesn't have free range in OS program memory**

How it limits the user:

- In the semantics:  $PC = (x8000 \mid UIMM8)$ 
  - What is the largest 8-bit number you can make? 256, **255** if you start from 0
  - e.g.:  $PC = x8000 \mid xFF = x80FF$

# Installing TRAPs Properly in OS Program Memory

```
; User Program Memory
```

```
.CODE
```

```
.ADDR x0000
```

```
; doing some fun stuff, like computing factorials!
```

```
; now, let's get a character from the keyboard!
```

```
;
```

```
TRAP x00
```

```
; OS Program Memory
```

```
.OS
```

```
.CODE
```

```
.ADDR x8000
```

```
SUB GETC
```

```
    LC R0, OS_KB_R_ADDR
```

```
    LDR R0, R0, #0
```

```
    BRzp GETC
```

```
    LC R0, OS_KB_R_ADDR
```

```
    LDR R0, R0, #0
```

```
RTI
```

- Because of TRAPs limitation,
  - *we shouldn't install our "TRAPS" starting at x8000*
- Why not?
  - *For one, user's might jump into the middle of our trap!*
  - *Imagine: TRAP x01? We'd jump right into LDR R0,...*
- Another reason?
  - Say our TRAP is > 256 lines long,
  - we couldn't have more than 1 TRAP

# Installing TRAPs Properly in OS Program Memory

Since TRAP instruction can only jump to the first 256 rows of OS Program Memory...

- We'll make the first 256 rows JMP instructions!
- They will serve as a jumping off point to TRAPS deeper in the OS Program Memory

```
.OS
.CODE
.ADDR x8000
    JMP TRAP_GETC          ; x00
    JMP TRAP_PUTC          ; x01
    JMP TRAP_DRAW_H_LINE; x02
    ...
    JMP TRAP_TURNOFF_CPU; xFF
```

## The first 256 lines of OS Program Memory

- called the: **TRAP VECTOR TABLE**
- We publish this list to the user:
  - user can call the TRAPS by number:
  - e.g.: TRAP x01, will call TRAP: PUTC
  - the table listing helps them map # to TRAP

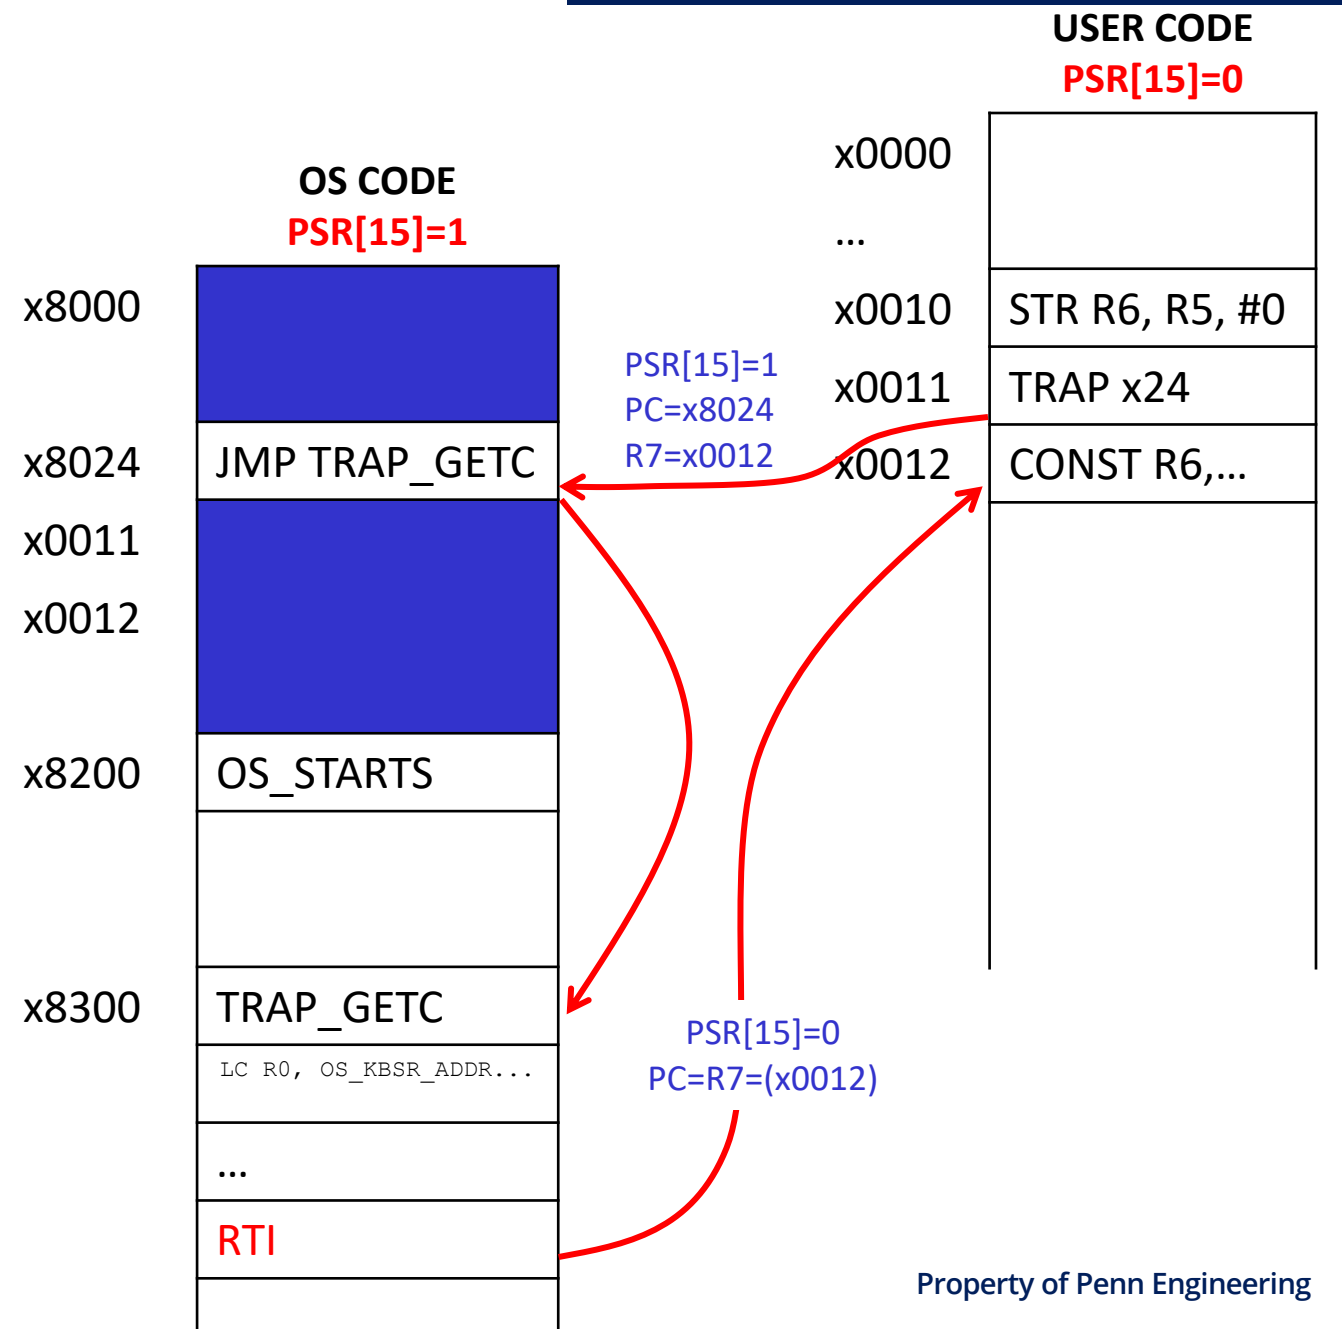
# Installing TRAPs Properly in OS Program Memory

```
.OS
.CODE
.ADDR x8300
TRAP_GETC                ;; this is TRAP x00
    LC R0, OS_KBSR_ADDR
    LDR R0, R0, #0
    BRzp TRAP_GETC
    LC R0, OS_KBDR_ADDR
    LDR R0, R0, #0
    RTI

TRAP_PUTC                ;; this is TRAP x01
    LC R1, OS_ADSR_ADDR
    LDR R1, R1, #0
    BRzp TRAP_PUTC
    LC R1, OS_ADDR_ADDR
    STR R0, R1, #0
    RTI
```

# Anatomy of a Trap

- When a TRAP is called:
  - CPU sets PSR[15]=1,
  - stores PC+1 in R7
  - and Jumps to entry in the TRAP Table
- This entry is another JMP instruction which redirects to the TRAP routine
- After the TRAP routine is complete:
  - it returns by using RTI,
  - which sets the PC to R7
  - which should contain the return address
  - and sets PSR[15] = 0



# A Word On What We've Created...

This system of TRAP calling, vector table organization, protection of the hardware

- This **system** of handling I/O...this way we are **operating**...
- Becomes known as the **Operating System** for the LC4
- The Operating System essentially serves as the means to handle I/O
- In the next segment, we'll abstract some of these ideas...
  - But I wanted you to know why we call this **OS** Program Memory!

# PASSING DATA TO/FROM TRAPS



# Things to Note About TRAP Calls:

## 1) The TRAP subroutine uses register file too!

- You may be using register file in your code in user program memory
  - The TRAP may overwrite it!
  - **Don't expect register file to be in the same shape as when you called the TRAP**

## 2) Two ways to pass arguments to TRAPS:

- Use a register in the register file

Or

- Use data memory and pass a starting address in data memory to the TRAP

# Backing Up the Register File

The register file will be used inside a TRAP

- *It will likely overwrite everything in the REGFILE*
- BEFORE you call a TRAP, save relevant content of REGFILE
- LDR and STR's "OFFSET" comes in handy here:

```

TEMPS .UCONST x4200      ; address of temporary storage
LC R7, TEMPS              ; load address into R7
STR R0, R7, #0             ; store R0 in TEMPS[0]
STR R1, R7, #1             ; store R1 in TEMPS[1]
STR R2, R7, #2             ; store R2 in TEMPS[2]
...
STR R6, R7, #6             ; store R6 in TEMPS[6]
TRAP x03                  ; call the TRAP
LC R7, TEMPS              ; load address into R7
LDR R0, R7, #0             ; restore R0 from TEMPS[0]
LDR R1, R7, #1             ; restore R1 from TEMPS[1]
LDR R2, R7, #2             ; restore R2 from TEMPS[2]

```

**Save content of REGFILE  
before you call TRAP**

**Restore content of REGFILE  
AFTER you return from TRAP**

# Examples of Passing Data to/from Traps

## TRAP\_GETC

- Uses register file's R0 to return data to the calling program

## TRAP\_PUTC

- Uses register file's R1 to get data from the calling program

## TRAP\_PUTS

- For your HW
- Before you call the trap, you place a "string" or ASCII characters in user data memory:
  - Ex: x4000 = 'T'  
x4001 = 'O'  
x4002 = 'M'
- Next, set register file's R0 = x4000 (starting address in data mem)
- Call the trap, passing in R0
- Trap does an LDR from R0, getting 'T' from user data mem
  - Continues looping loading characters until data = 0

# OPERATING SYSTEMS – LOGICAL DISCUSSION

# Operating Systems (OSes)

## First job of an OS:

- Handle I/O
- OSes virtualize the hardware for user applications

## In real systems, only the operating system (OS) does I/O

- “User” programs *ask* OS to perform I/O on their behalf
- 3 reasons for this setup:

### 1) Standardization

- I/O device interfaces are nasty, and there are many of them
- Think of disk interfaces: S-ATA, iSCSI, IDE
- User programs shouldn't have to deal with these interfaces
  - In fact, even OS doesn't have to deal with most of them
  - Most are buried in “device drivers”

# Operating Systems (OSes)

## 2) Raise the level of **abstraction**

- Wrap nasty physical interfaces with nice logical ones
  - Wrap disk layout in file system interface

## 3) Enforce **isolation** (usually with help from hardware)

- Each user program thinks it has the hardware to itself
  - User programs unaware of other programs or (mostly) OS
- Makes programs much easier to write
- Makes the whole system more stable and secure
  - A can't mess with B if it doesn't even know B exists

# In the Real World

- This system call/TRAP mechanism is commonly used in actual systems
- As an example the BIOS (Basic Input Output System) on many Intel PCs provided precisely this functionality to allow programs to access basic input and output devices, keyboards, displays, timers etc.
- More modern systems use EFI (Extensible Firmware Interface) which is a more sophisticated version of the same thing.

# WHAT IS SYNCHRONIZATION



# What Is Synchronization?

## Two schemes for interacting with I/O devices:

### 1) Way have seen so far is **polling**

- “Are we there yet? Are we there yet? Are we there yet?”
- CPU keeps checking status register in a loop
- Very inefficient, CPU has better things to do

### 2) Alternative scheme is called **interrupts**

- “Wake me when we get there.”
- Device sends special signal to CPU when status changes
- CPU stops current program, saves its state
- CPU “handles interrupt”: checks status, moves data
- CPU resumes stopped program, as if nothing happened!!!!!!!!!!

# LC4 Simulator Does Not Support Interrupts



## Although LC-3 ISA does

- Presumably LC4 does too, but haven't tried
- And could be tricky without PC-relative LD/ST

## Why not?

- Suppose program wants to read from the keyboard
- Interrupts make sense if CPU could do something else meantime
- Program itself can't do anything, it's waiting for the key ...
- ...so another program has to run while it's waiting
- And LC4 doesn't support multi-programming
- So there is no real point ☺

# How Interrupts Work



## **CPU in user mode**

Doing boring user work  
partying etc. etc.

## **Holy smokes**

A device is ready for input or output  
Better raise the interrupt!

## **CPU switches to superhero mode:**

Lays a smackdown on the I/O device.  
When done return to regular life  
partying like nothing happened!

# What Is Synchronization?

## **Interrupt driven I/O is **asynchronous****

- Highly overloaded word, here means “takes unbounded time”

## **Asynchronous programming is difficult to say the least!**

- Suspend, find something else to do, resume on signal

## **Better leave this to the OS**

- Suspends you transparently
- Runs another program while you wait (it knows about these)
- Resumes you transparently on signal

## **To you, I/O looks **synchronous****

- Similarly overloaded, here means “happens immediately”